

Wyatt Walsh

SENIOR AI/ML ENGINEER

wyattowalsh@gmail.com (209) 602-2545 New York City, New York [linkedin.com/in/wyattowalsh](https://www.linkedin.com/in/wyattowalsh) github.com/wyattowalsh

Summary: Senior AI/ML engineer building agentic AI, LLM document intelligence, data pipelines, and developer tooling for fintech, risk, and compliance teams.

EXPERIENCE

Senior AI/ML Engineer

Frame Payments | Remote | Oct 2025 - Jan 2026

- Deployed production LLM document intelligence with PydanticAI, Amazon Bedrock AgentCore, and Claude; added human-in-the-loop browser feedback for confidential uploads and reduced analyst follow-ups/resubmissions ~90% by blocking miscategorized documents before review.
- Built multimodal restricted-item compliance screening by combining product images and structured metadata; translated banking-partner policy requirements into automated risk checks for merchant onboarding across tens of merchants and multi-location catalogs.

Associate III, Senior Software Engineer - Data Integrations

JPMorgan Chase & Co. | New York, NY | May 2023 - Apr 2025

- Pioneered firm-wide rollout of secure generative AI developer tools (LLMs, GitHub Copilot, Codeium) as one of the first 500 engineers, shaping scalable deployment and leading workshops for technical and non-technical teams.
- Devised ML-driven data-quality anomaly detection for credit trading and risk analytics, reducing false positives by 25% and cutting manual review time so risk teams could focus on critical alerts.
- Developed resilient Python wrapper for KDB+, democratizing data access for 5,000+ users with CI-style tests and optimized Q functions, reducing manual processing by over 10 hours per week.

Senior Software Engineer

Stealth Web3 Startup | Remote | Mar 2022 - Jul 2022

- Architected a scalable NFT-collateralized lending DApp using React, Next.js, TypeScript, MobX, and ethers.js, enabling secure, performant on-chain loans.
- Implemented end-to-end CI/CD with GitHub Actions, Husky, Cypress, and Jest to halve release cycles and raise test coverage to 90%, while mentoring a front-end intern to build key UI components.

Founder, Chief Technology Officer

SandLabs | Remote | Jun 2021 - Mar 2022

- Engineered and launched a Next.js/TailwindCSS marketing site with top-tier Lighthouse scores and secured \$75K in OCEAN tokens via a competitive DAO grant.
- Engineered Python pipelines to extract and normalize a 1M+ Web3-related Reddit posts & comments dataset and published the data on Ocean Protocol's Data Marketplace.

Demand Forecasting Extern

Gap Inc. | San Francisco, CA | Jan 2020 - May 2020

- Applied time-series decomposition and bias correction techniques to demand forecasts, boosting accuracy by up to 25% and minimizing surplus inventory costs.
- Owned the data engineering layer by architecting scalable MySQL and Apache Spark pipelines across 5+ years of sales data, improving data processing throughput by 95% and standardizing insights across forecasting teams.

SKILLS

MACHINE LEARNING AND GENERATIVE AI: scikit-learn, Hugging Face Transformers, OpenAI API, LangChain, LangGraph, Model Context Protocol (MCP), Retrieval Augmented Generation (RAG), Prompt Engineering, PydanticAI, Amazon Bedrock, Claude, GitHub Copilot

DATA ENGINEERING AND ANALYTICS: Pandas, NumPy, SQLAlchemy, Pydantic, Pandera, Apache Spark, Amazon Kinesis, AWS Glue, KDB+, PostgreSQL, DynamoDB, ETL, SQLite

PROGRAMMING & SCRIPTING LANGUAGES: Python, SQL, TypeScript, JavaScript, Go, Bash

CLOUD INFRASTRUCTURE AND DEVOPS: AWS, Docker, Kubernetes, Terraform, GitHub Actions, Serverless Architecture, AWS Lambda, ECS, EKS, Vercel

PROJECTS

AI Agent Harness Configs - LLM workflow catalog with reusable skills, MCP-ready configs, and cross-runtime agent automation.

NBA Basketball Database - Validated data warehouse spanning 65K+ NBA games and 13M+ events with ETL, docs, CLI, and Kaggle distribution.

ProxyWhirl - Data-collection toolkit with nine proxy rotation strategies, 100+ sources, and self-healing health checks.

AI/ML Web Feeds - Typed AI/ML source-discovery platform with validation and docs-backed feed workflows for model and agent research.

MCP-Crawl4AI - FastMCP server that lets AI agents crawl, extract, and return structured web content over stdio/HTTP.

EDUCATION

Bachelor of Science

University of California, Berkeley - College of Engineering, Industrial Engineering & Operations Research
Aug 2014 - Dec 2020

High School Diploma

The Hotchkiss School
Aug 2012 - May 2014

CERTIFICATIONS

AWS Certified Cloud Practitioner | IBM Data Science Professional Certificate